



BIG DATA CASE STUDY NIGHT #3



STREAM AND ANALYZE TIME SERIES IN REAL-TIME WITH AWS

John Mousa

Sr. Solutions Architect
Amazon Web Services

John Mousa



Engineer at heart, enabling enterprise customers in DACH and world wide bridge the worlds of microservices and analytics as a Sr. Solutions Architect at Amazon Web Services.



/in/johnmousa

/johnmousa

@JohnMousa_



Agenda for Today



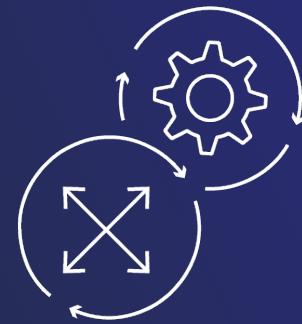
Our use case for
Power and Utilities



Overview on
Architecture



Solutions

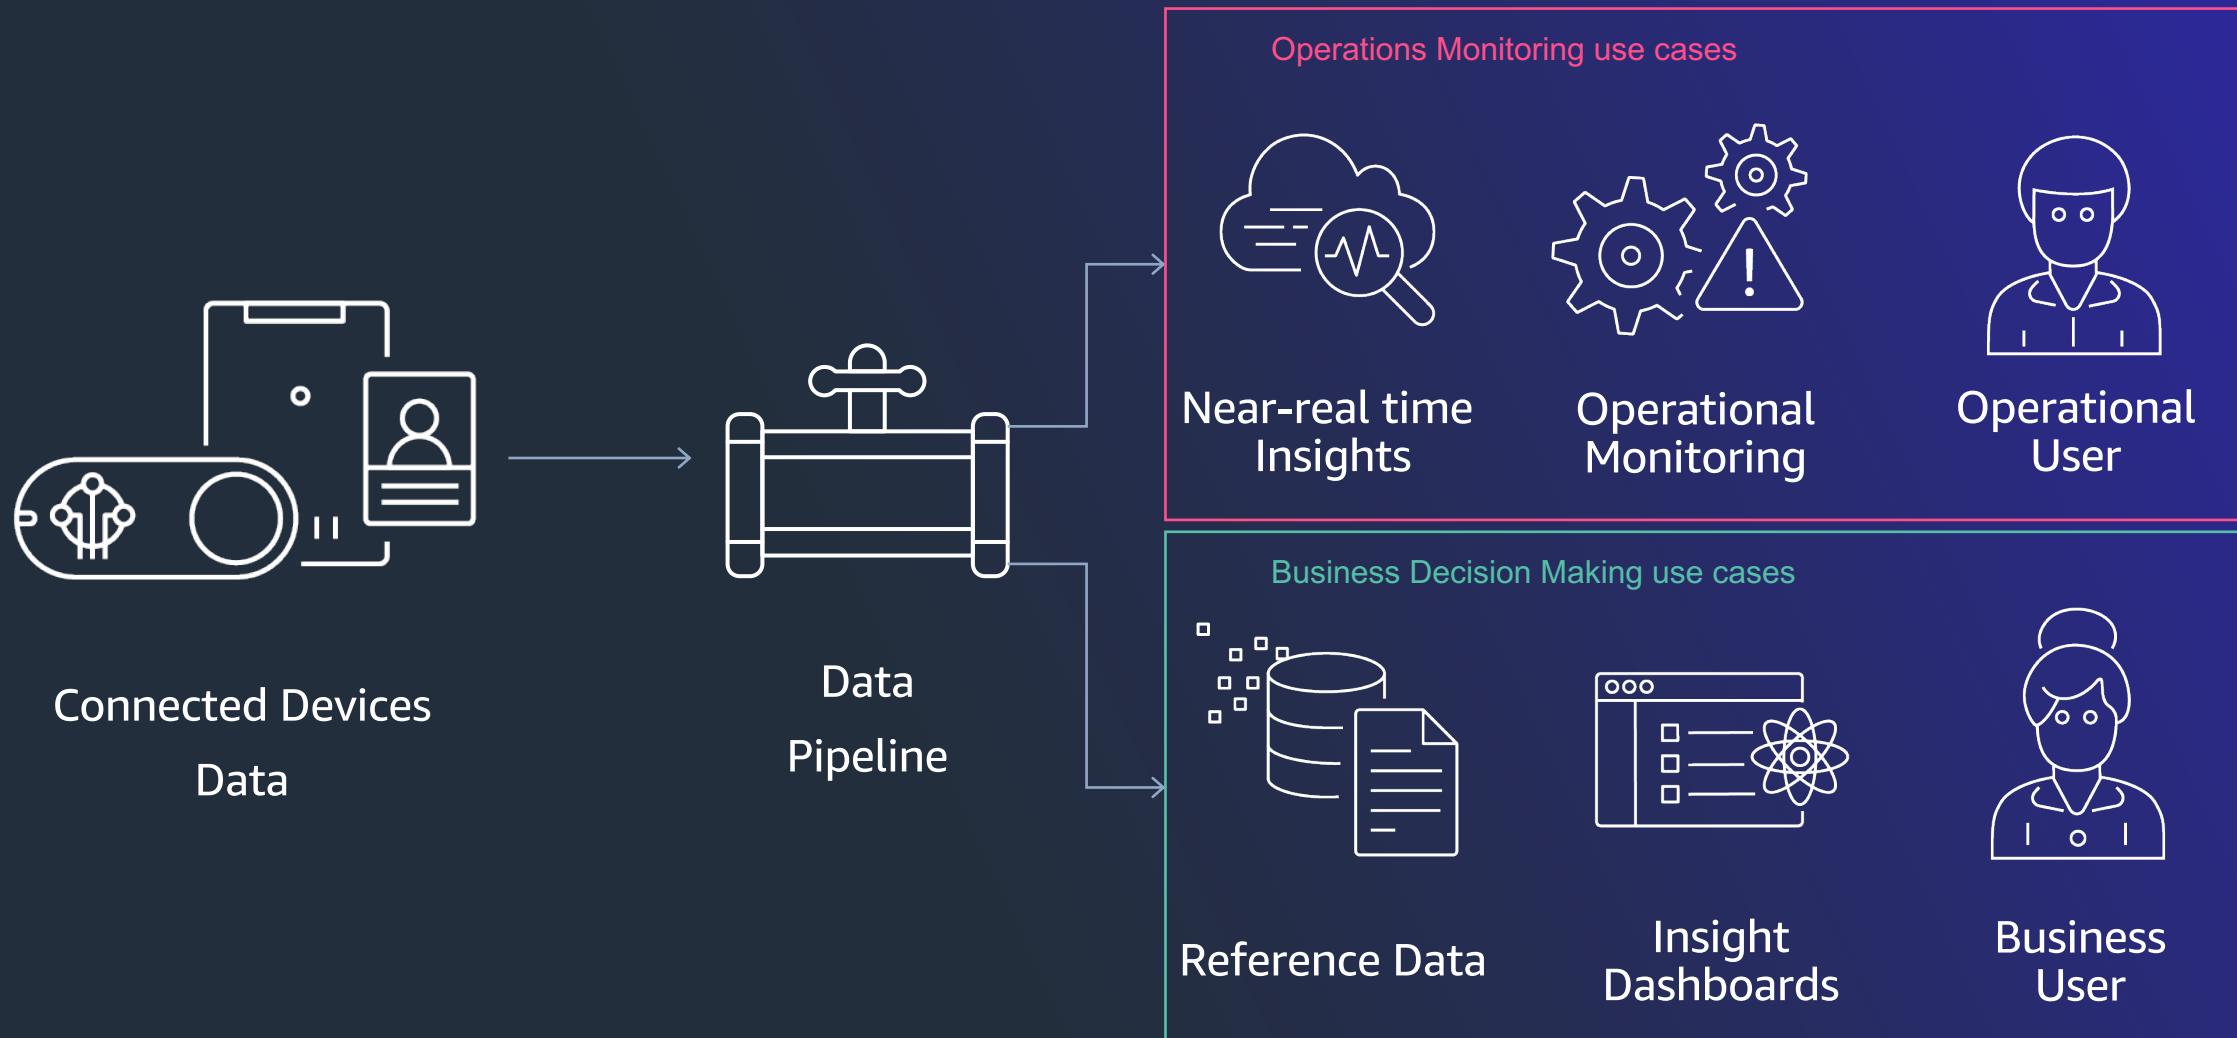


Operational View



Strategic View

Collecting Telemetry Data from Connected Devices



Architectures

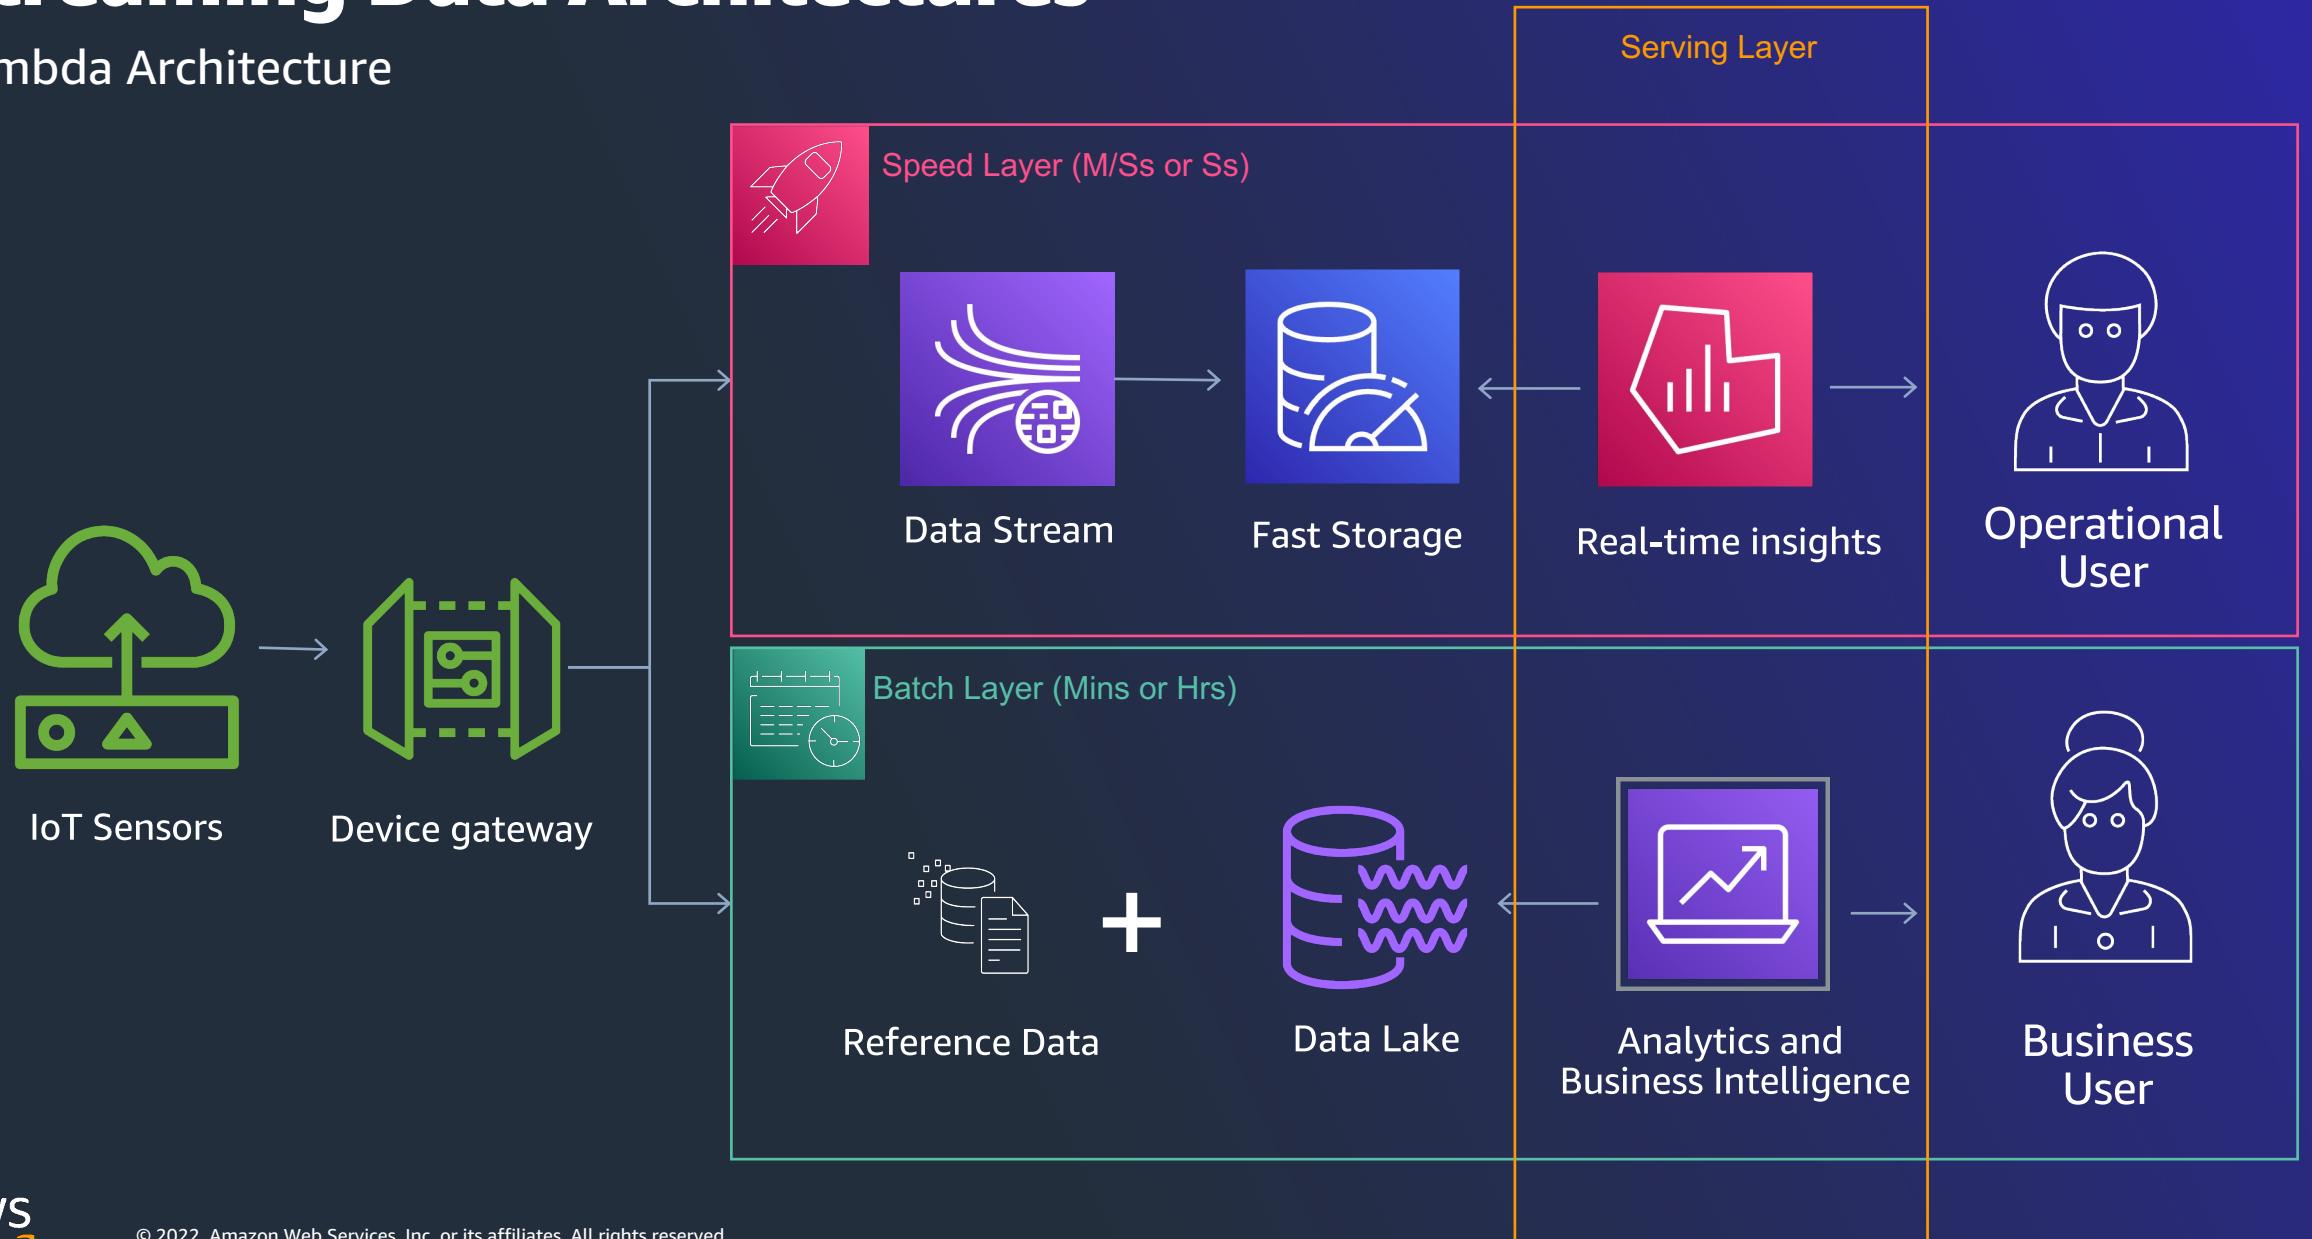


© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.



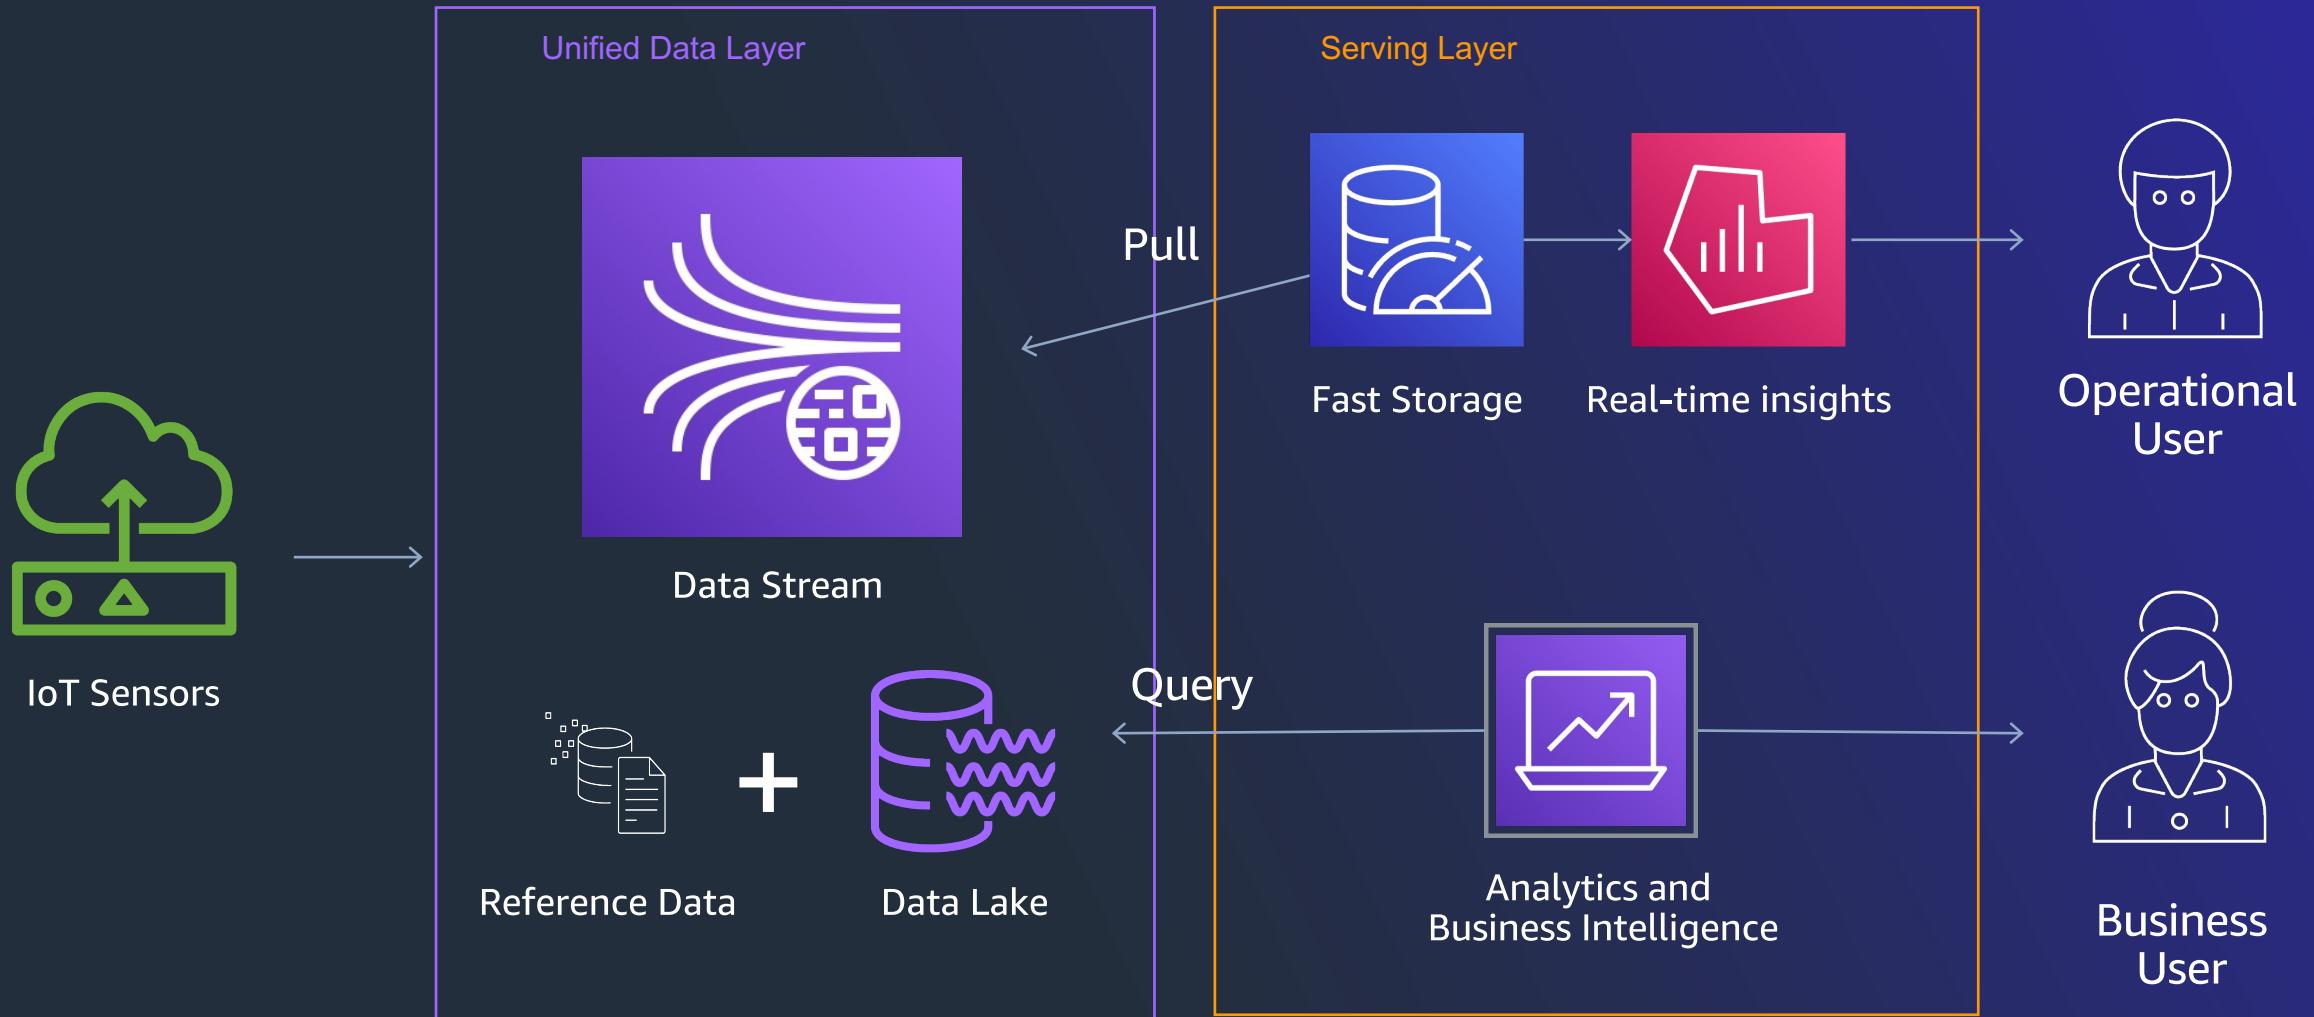
Streaming Data Architectures

Lambda Architecture



Streaming Data Architectures

Kappa Architecture



Streaming Data Architectures

Lambda VS Kappa

λ

- Batch layer is built for big data
- Result in overhead due to code and resource duplication

K

- No redundancies
- Batch workloads needs to be in a sink

Solutions



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.



But Why Did We Consider Such Solutions



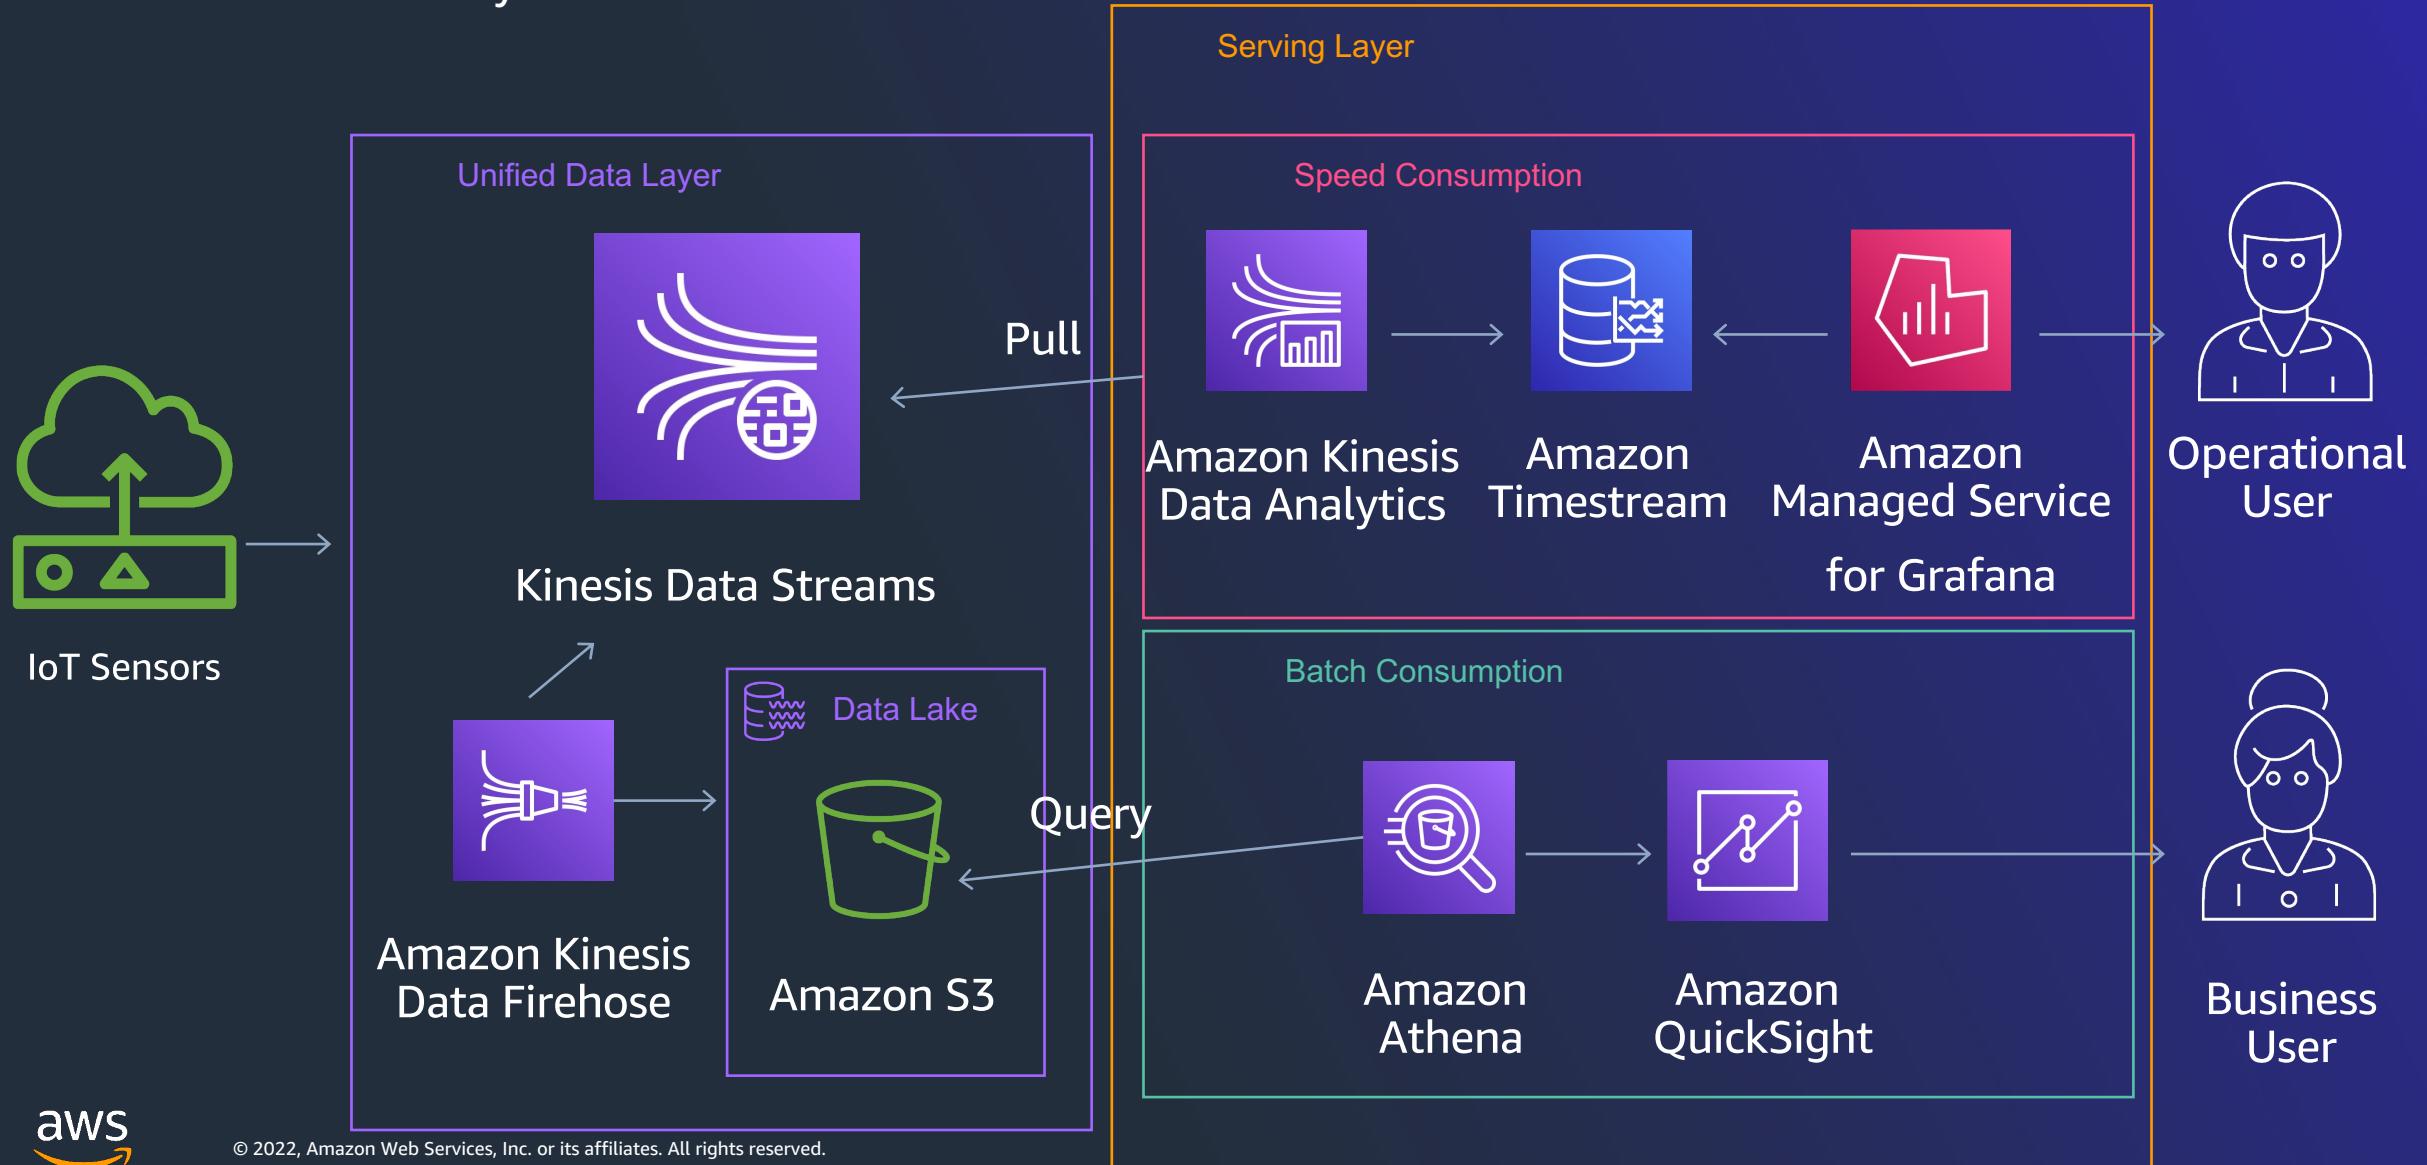
Streaming is Expensive

Streaming is Lossy

Hard to Operate,
Scale and Maintain

Current Data Architecture

Enter Serverless Analytics



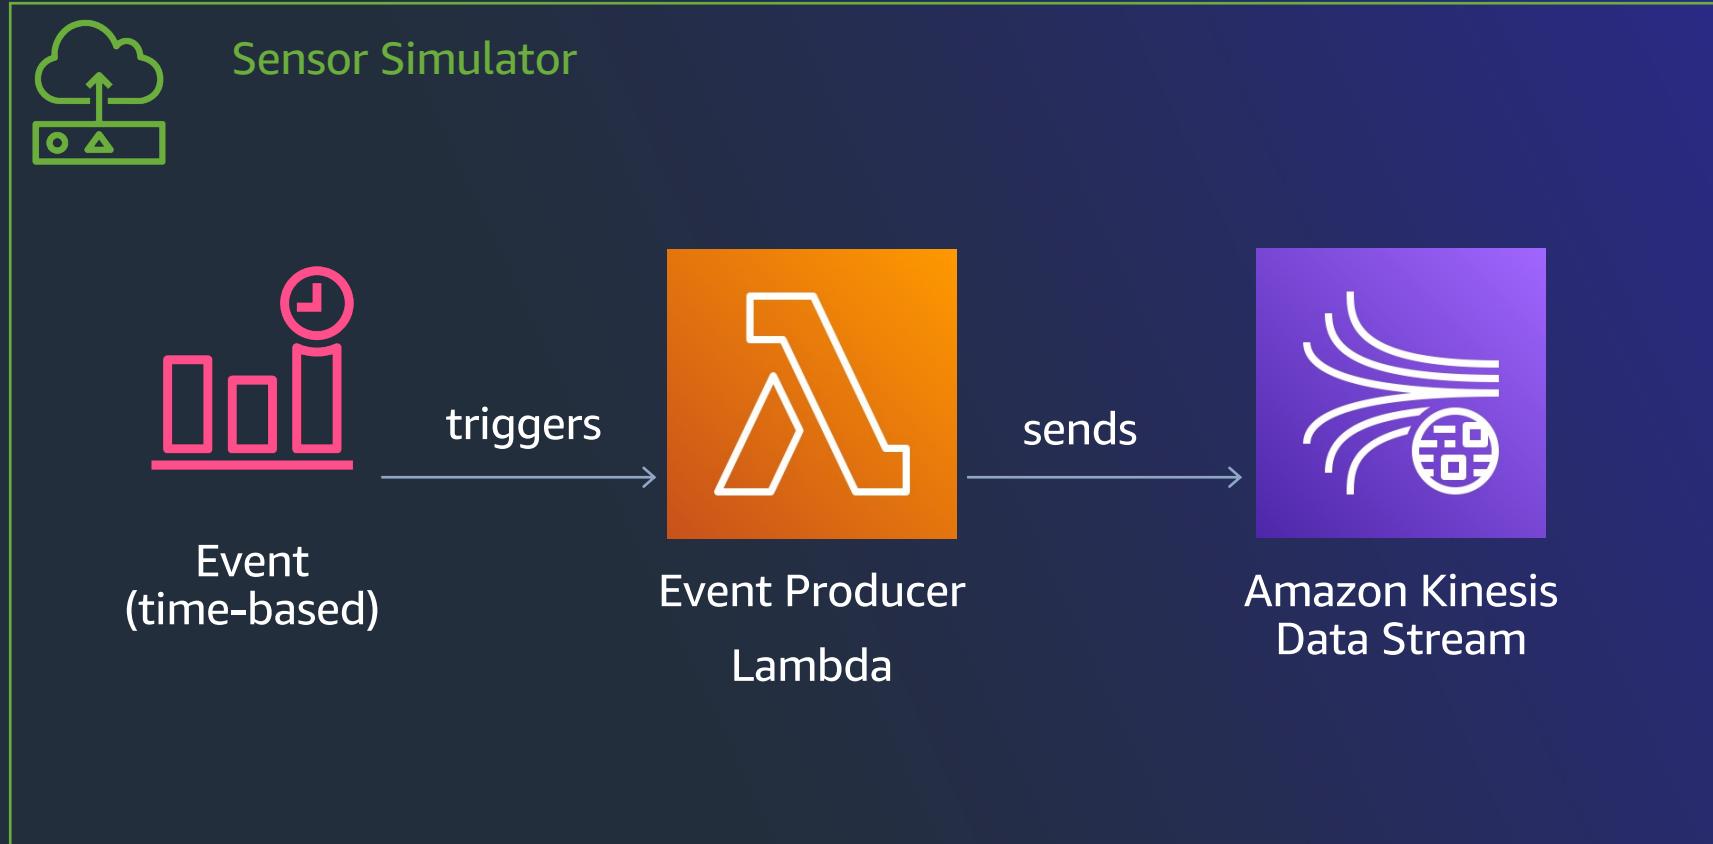
Ingestion



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Data Ingestion from Devices



Data Ingestion from Devices

CDK Infrastructure

```
sample_device_producer = aws_lambda_python.PythonFunction(self, 'SampleDeviceProducer',
    entry='stacks/sample_kinesis_stream_producer/producer_lambda',
    index='app.py',
    runtime=aws_lambda.Runtime.PYTHON_3_8,
    timeout=core.Duration.seconds(30))

stream.grant_write(sample_device_producer)

lambda_input = {"Stream": stream.stream_name}
Rule(self, 'ProducerTriggerEventRule',
    enabled=True,
    schedule=Schedule.rate(Duration.minutes(1)),
    targets=[aws_events_targets.LambdaFunction(handler=sample_device_producer,
        event=RuleTargetInput.from_object(lambda_input))])
```



Data Ingestion from Devices

Lambda Writing to Kinesis Data Streams

```
def write_records(stream_name, records):
    kinesis_records = []
    for record in records:
        kinesis_records.append({
            'Data': json.dumps(record),
            'PartitionKey': record["DeviceID"]
        })

    result = None
    try:
        result = kinesis.put_records(
            StreamName=stream_name,
            Records=kinesis_records,
        )

        status = result['ResponseMetadata']['HTTPStatusCode']
        print("Processed {} records. WriteRecords Status: {}".format(len(records), status))
    except Exception as err:
        print("Error:", err)
        if result is not None:
            print("Result:{}".format(result))
```

```
def handler(event, context):
    sent = 0
    for x in range(iterations):
        records = []
        for device_id in device_ids:
            records.append(prepare_record(measures, device_id,
                                            (iterations - x) / iterations * 60))

        print("records {}".format(len(records)))
        write_records(event["Stream"], records)
        sent = sent + len(records)
    return {"records": sent}
```

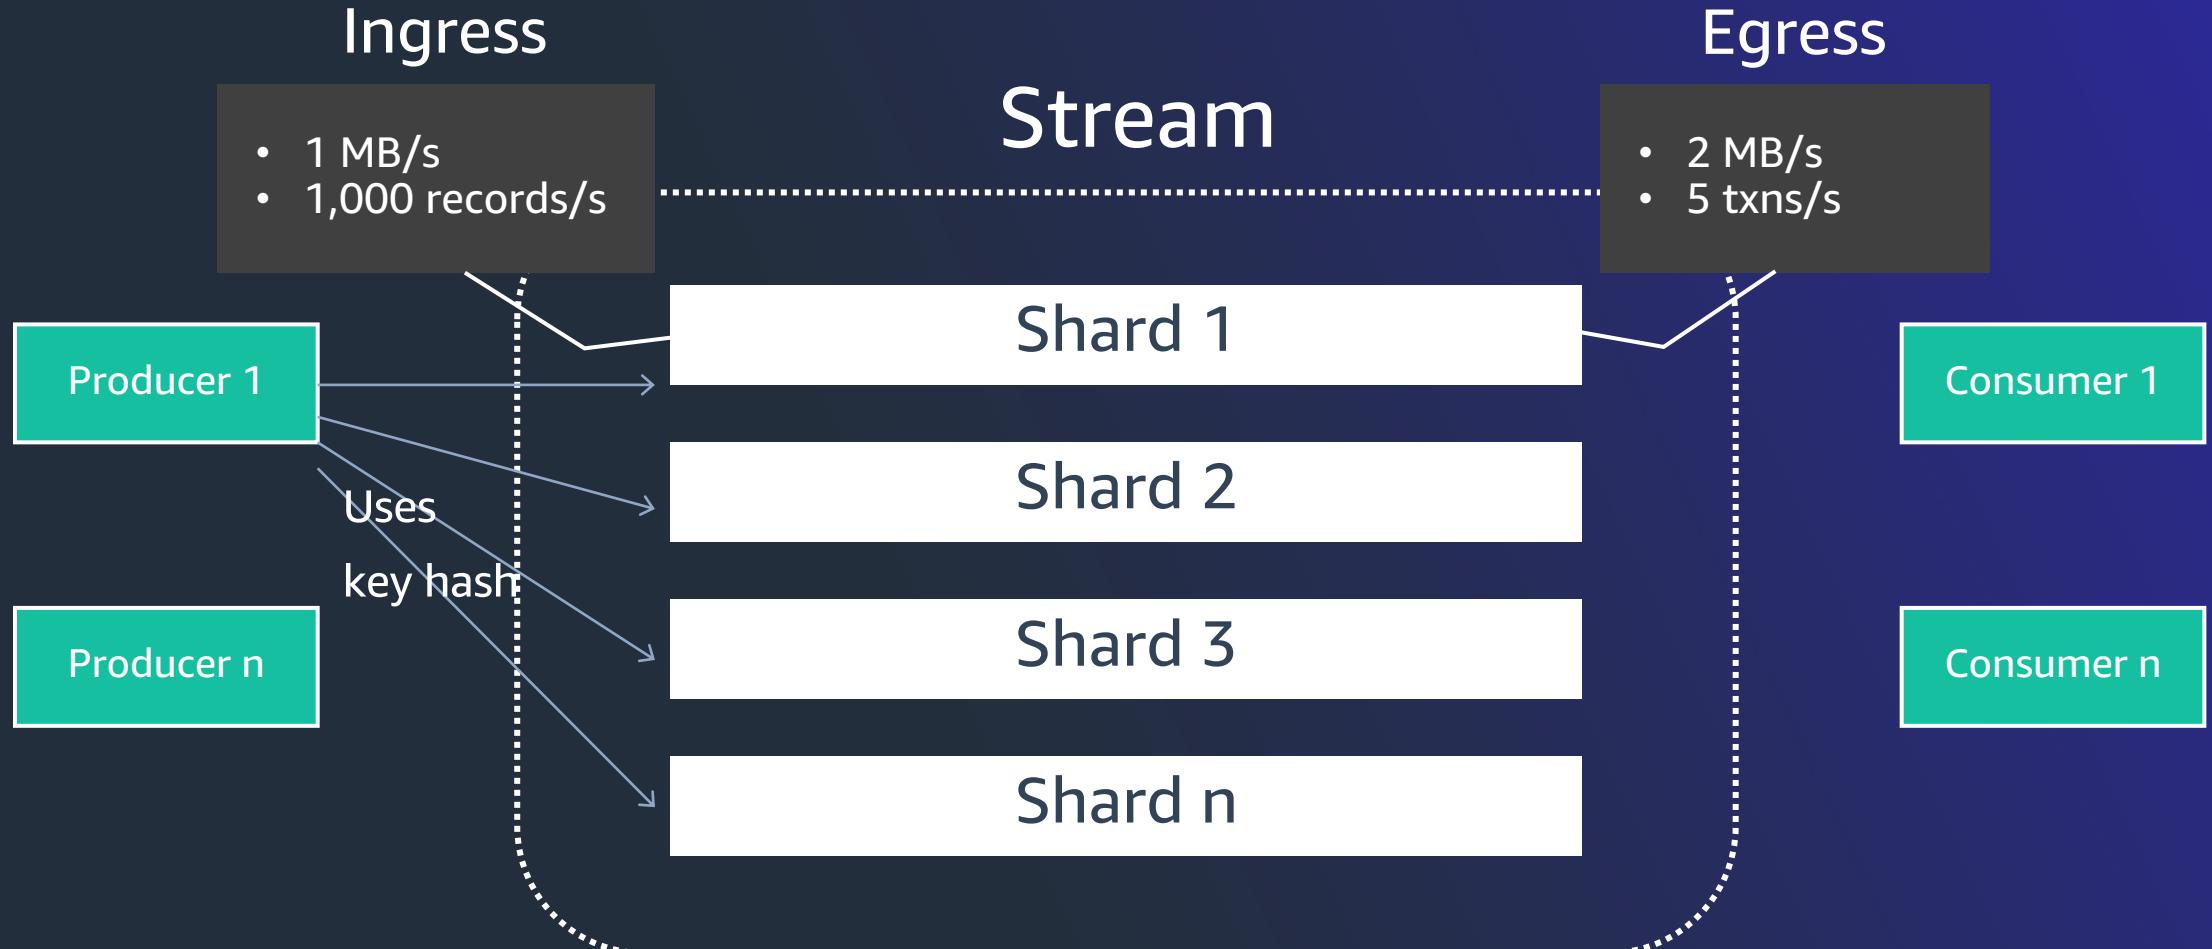


Data Streaming



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

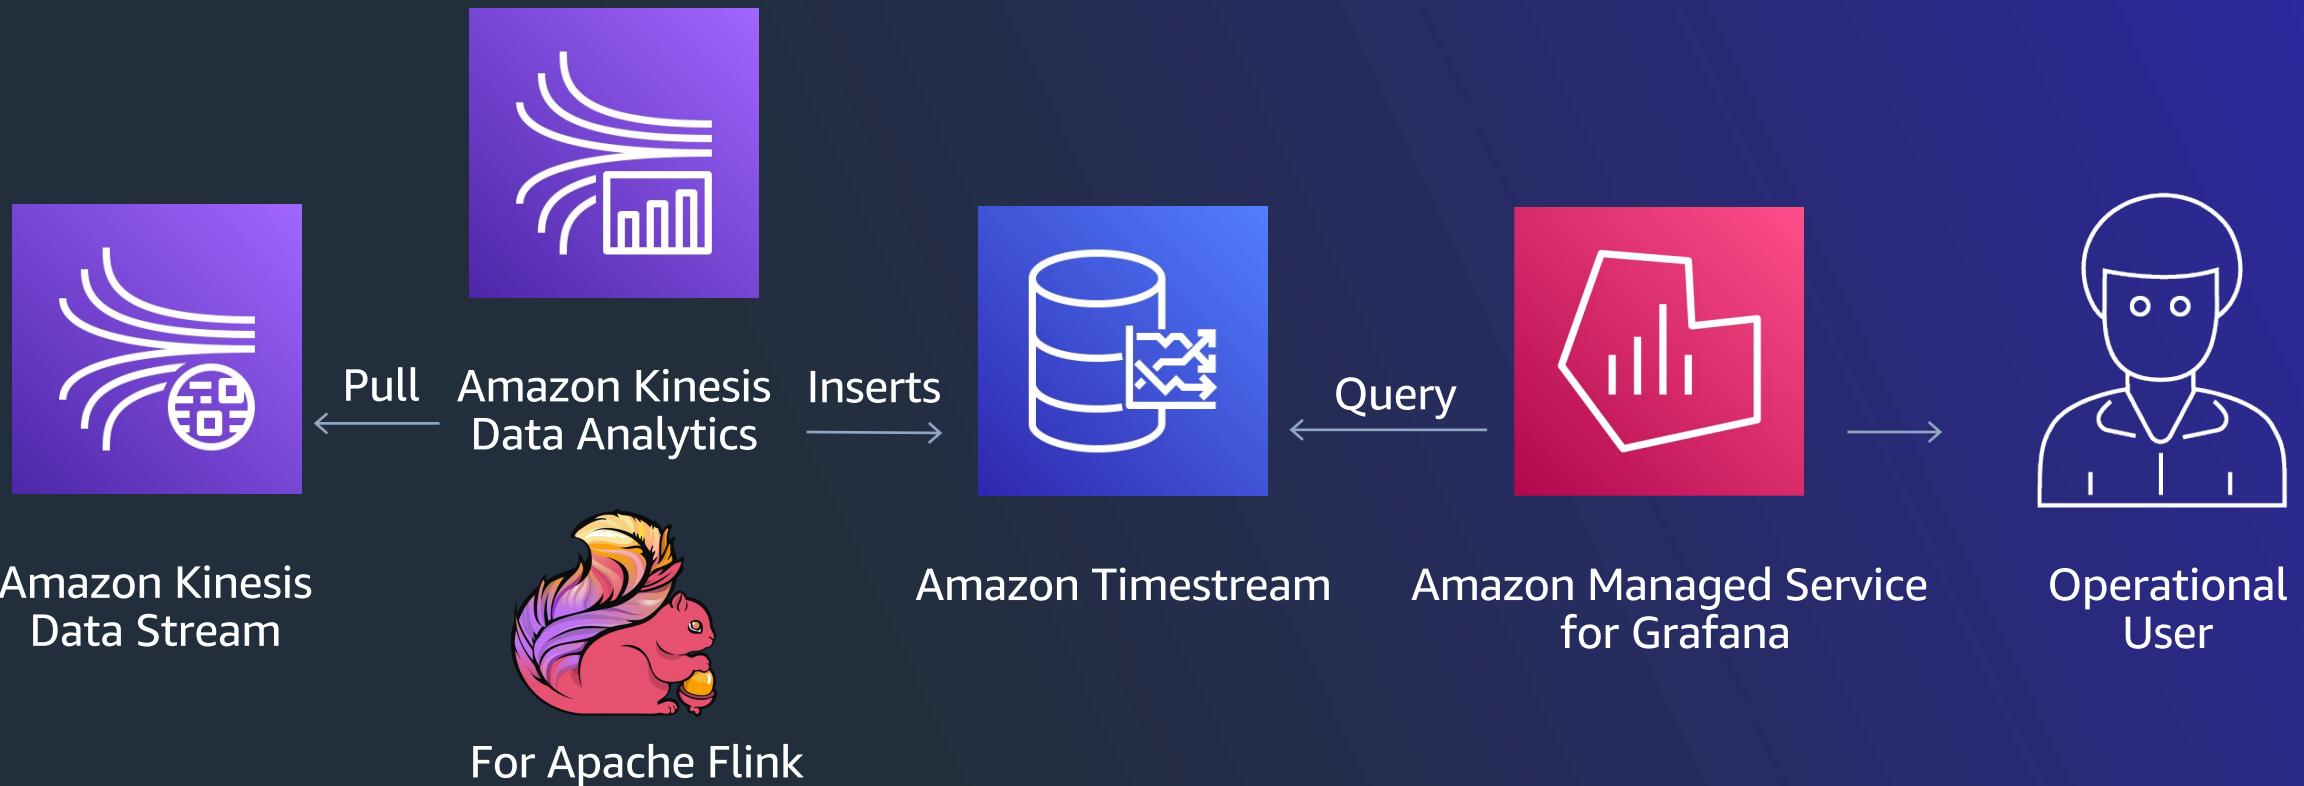
Kinesis Data Streams: Granular Provisioning and Scaling



Operational View

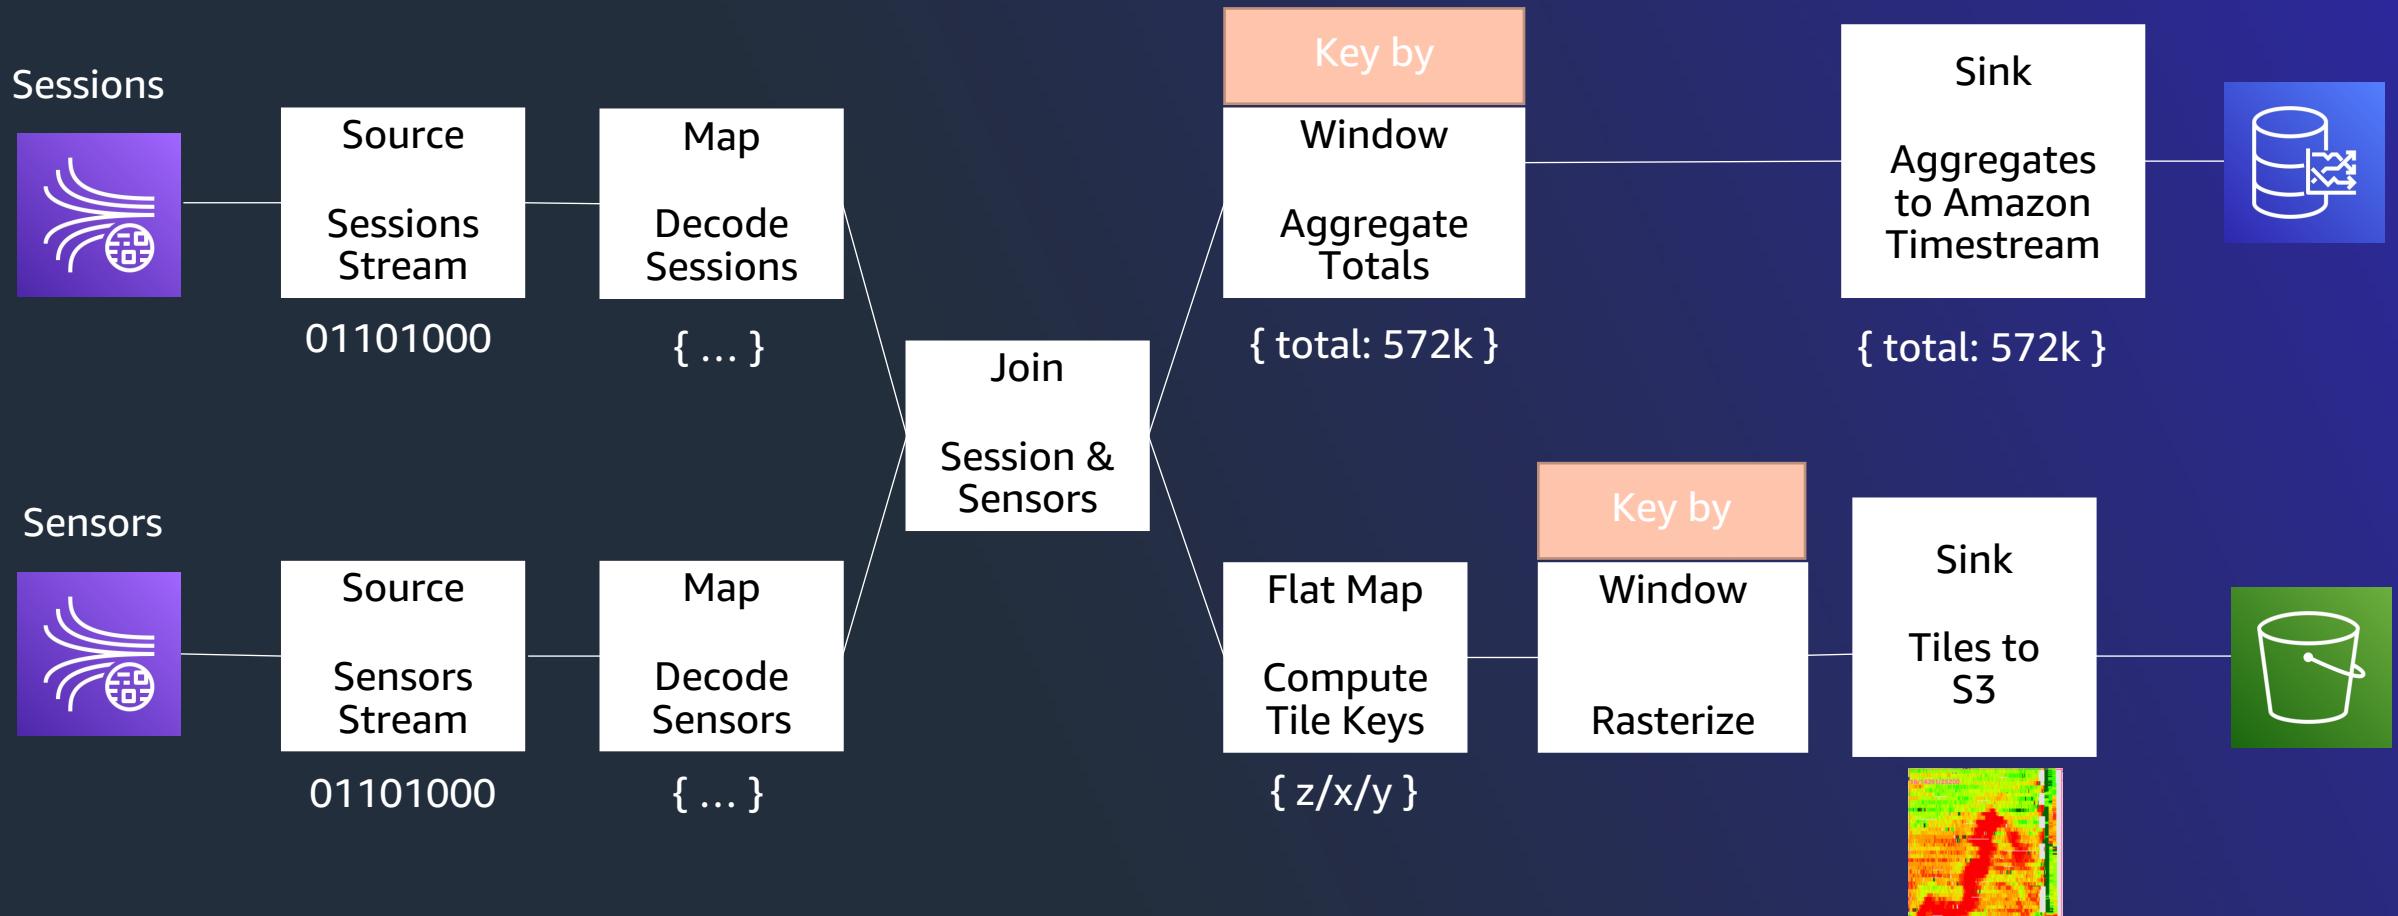


Near-Real Time Operational Insights Architecture



Streaming Analytics

Apache Flink - Example



Streaming Analytics

Apache Flink – Declarative Streams

```
createKinesisSource(env, parameter)
    .map(JsonToTimestreamPayloadFn()).name("MaptoTimestreamPayload")
    .process(OffsetFutureTimestreamPoints()).name("UpdateFutureOffsetedTimestreamPoints")
    .addSink(TimestreamSink(region, databaseName, tableName, batchSize))
    .name("TimestreamSink<$databaseName, $tableName>")
```



Amazon Timestream Architecture

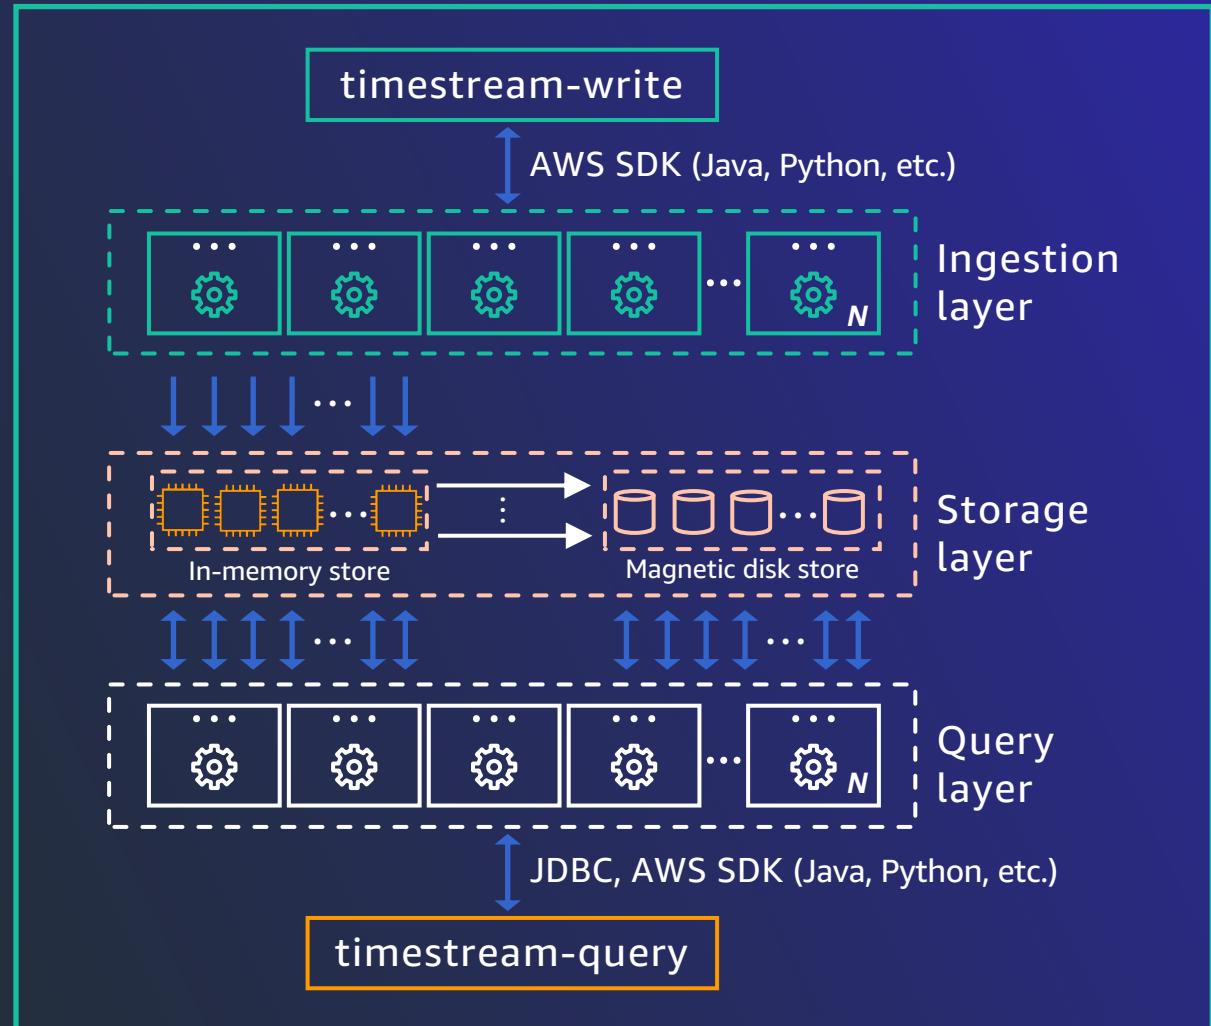
Purpose-built for time series workloads

Decoupled architecture

- Highly available
- **Independently** scalable ingestion, storage, and query processing

High throughput auto scaling ingestion

- **Durable**: data is **replicated** across multiple Availability Zones
- Automatic data deduplication handling
- No need to provision or configure resources



Amazon Timestream Architecture

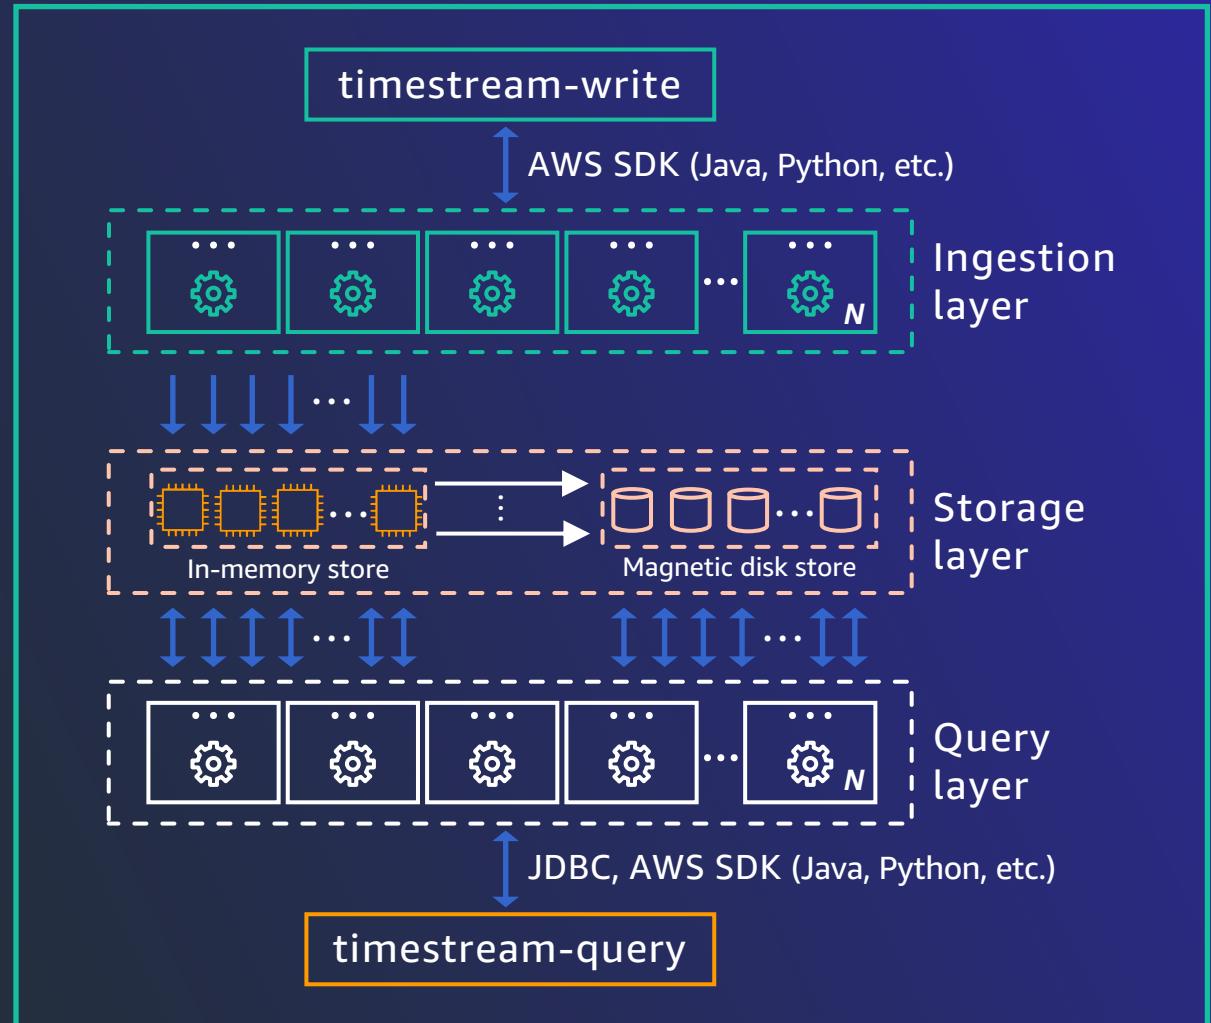
Serverless, cloud-native AWS built technology

Multiple tiers of storage

- Scalable to petabytes and beyond
- **Memory store**: fast point-in-time queries, high-throughput durable ingestion
- **Magnetic store**: high-performance analytical queries and low-cost long-term storage

Scalable SQL query engine

- **Adaptive query engine** is capable of querying data across multiple data tiers
- **No indexes** to configure and **no provisioning** required



Operational Views

Amazon Managed Service for Grafana



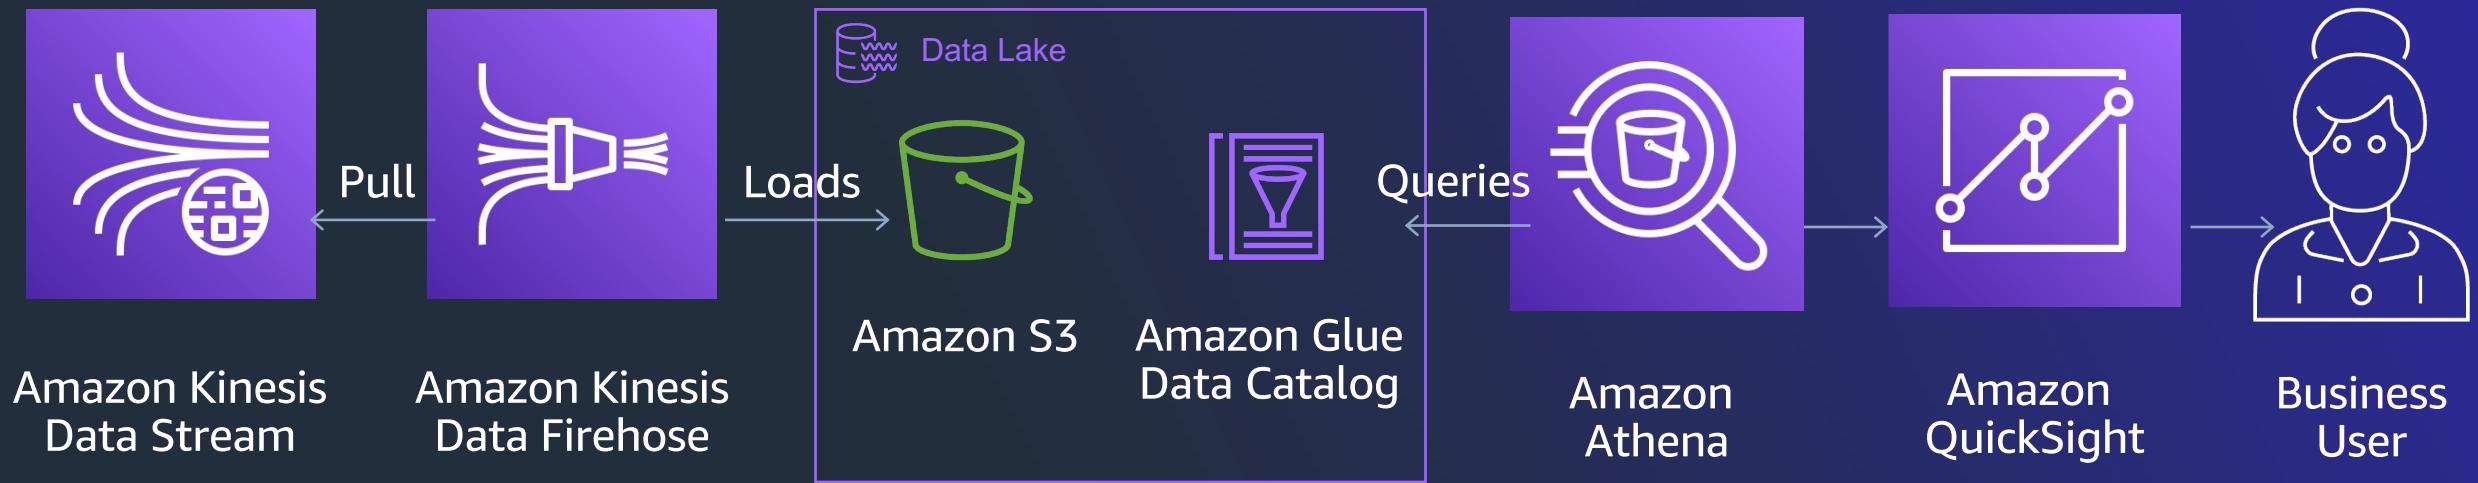
© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Strategic Big Data Insights



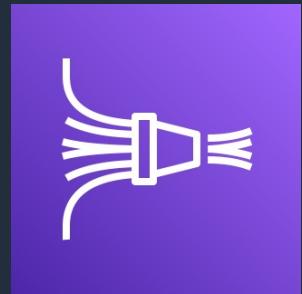
© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Strategic Big Data Insights Architecture



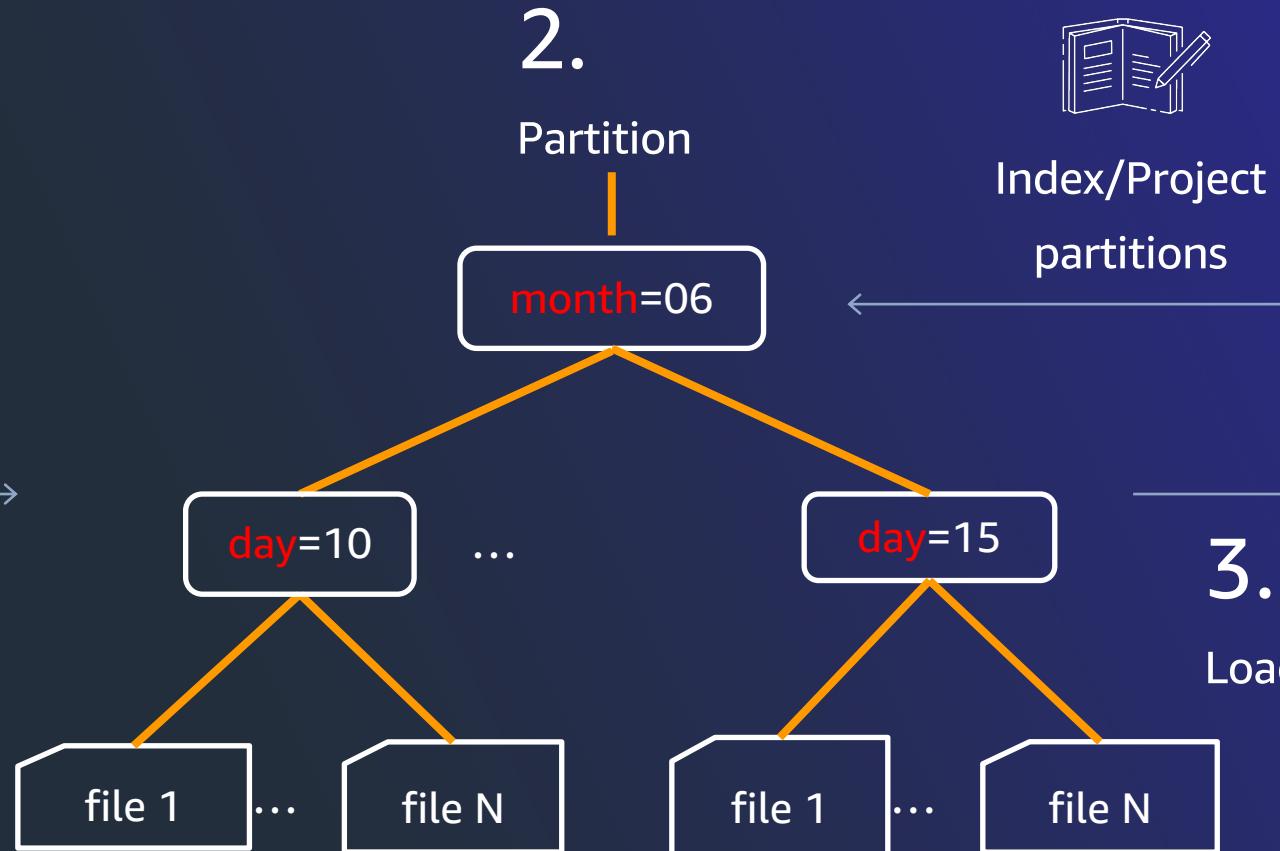
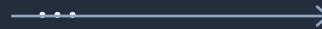
Data Encoding and Partitioning

Amazon Kinesis Data Firehose



Amazon Kinesis
Data Firehose

1.
 {JSON Lines}
 {Formatted}



Index/Project
partitions



Amazon Glue
Data Catalog

3.
Load

Amazon S3

Data Encoding and Partitioning

Amazon Kinesis Data Firehose



Discover Dynamic Partitions

Glue Data Catalog Partition Projection

```
Prefix: "readings/json/year=!{partitionKeyFromQuery:year}/month=!{partitionKeyFromQuery:month}  
/day=!{partitionKeyFromQuery:day}/hour=!{partitionKeyFromQuery:hour}"/
```

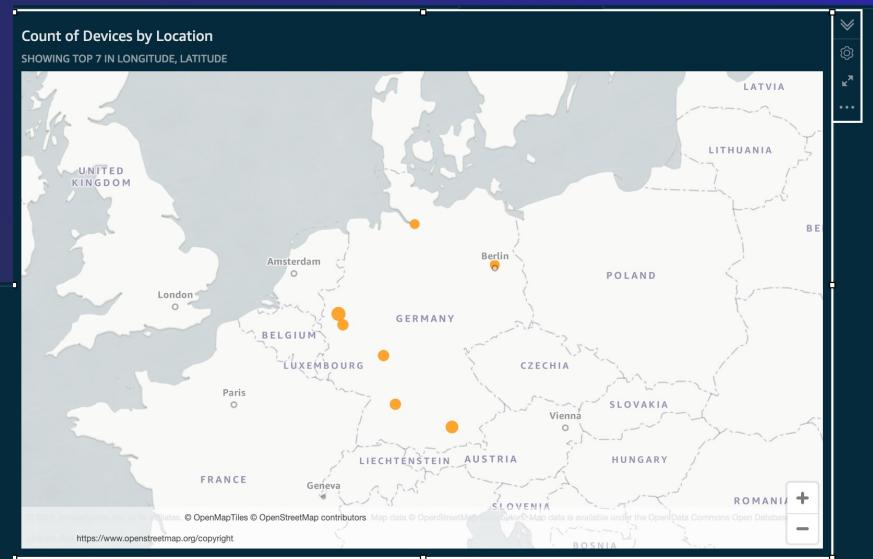
bucket_name > measures/ > json/ > year=2023/ > month=01/ > day=02/ > hour=01/

```
a  
CREATE EXTERNAL TABLE raw_device_data (  
  ...  
)  
...  
PARTITIONED BY (  
  year INT,  
  month INT,  
  day INT,  
  hour INT,  
)  
LOCATION "s3://DOC-EXAMPLE-BUCKET/prefix/"  
TBLPROPERTIES (  
  "projection.enabled" = "true",  
  "projection.year.type" = "integer",  
  "projection.year.range" = "2011,2099",  
  "projection.year.digits" = "4",  
  "projection.month.type" = "integer",  
  "projection.month.range" = "1,12",  
  "projection.month.digits" = "2",  
  "projection.day.type" = "integer",  
  "projection.day.range" = "1,31",  
  "projection.day.digits" = "2",  
  "projection.hour.type" = "integer",  
  "projection.hour.range" = "0,23",  
  "projection.hour.digits" = "2",  
  "storage.location.template" = "s3://DOC-EXAMPLE-BUCKET/prefix/year=${year}/month=${month}/day=${day}  
/hour=${hour}"/  
)
```

```
"projection.enabled" = "true",  
"projection.year.type" = "integer",  
"projection.year.range" = "2011,2099",  
"projection.year.digits" = "4",
```

Insights Views

Amazon QuickSight



Cost Factors

Dynamic (per load)

- Athena when queried / GB scanned
- Kinesis
 - Data ingested, per GB (Includes 24-hour retention) \$0.08
 - Data retrievals, per GB \$0.04
- Kinesis Firehose
- Amazon Timestream
 - 1 million write of 1KB size \$0.50

Static (even when no events come)

- Amazon S3 Storage
- Amazon Timestream
 - Price per GB stored per hour \$0.036
- Amazon Kinesis
 - Per stream, per hour \$0.04
 - Kinesis Data Analytics
 - Kinesis Processing Unit, Per Hour \$0.11 per hour
 - Running Application Storage, Per GB-month \$0.10 per GB-month

* Cost as of 28th of September 2022, in US-EAST-1 region, you pay for what you actually consume, checkout <https://calculator.aws>

Purpose-Built AWS Data Analytics Services for Your Modern Data Architecture

- Work backwards from context and use cases, don't build technical platforms
- Use the right tool for your use case
- AWS has multiple serverless analytics services for different needs
- You're one Experiment away



More Resources

- Near real-time processing with Amazon Kinesis, Amazon Timestream, and Grafana [Blog Post](#)
- Companion samples [github repository](#)
- Patterns for IoT time series ingestion [Blog Post](#)
- Smart meter analytics [QuickStart](#)



Please complete the session survey

Direct link to slides right after 5-click survey feedback

<https://eventbox.dev/survey/SQ1X96F>



Thank you!

John Mousa

Amazon Web Services
Sr. Solutions Architect | AWS DACH
Enterprise

